

Mean-Variance Optimization and the CAPM

These lecture notes provide an introduction to mean-variance analysis and the capital asset pricing model (CAPM). We begin with the mean-variance analysis of Markowitz (1952) when there is no risk-free asset and then move on to the case where there is a risk-free asset available. We also discuss the difficulties of implementing mean-variance analysis in practice and outline some approaches for resolving these difficulties. Because optimal asset allocations are typically very sensitive to estimates of expected returns and covariances, these approaches typically involve superior or more robust parameter estimation methods.

Mean-variance analysis leads directly to the *capital asset pricing model* or CAPM. The CAPM is a one-period equilibrium model that provides many important insights to the problem of asset pricing. The language / jargon associated with the CAPM has become ubiquitous in finance.

1 Markowitz's Mean-Variance Analysis

Consider a one-period market with n securities which have identical expected returns and variances, i.e. $E[R_i] = \mu$ and $\text{Var}(R_i) = \sigma^2$ for $i = 1, \dots, n$. We also suppose $\text{Cov}(R_i, R_j) = 0$ for all $i \neq j$. Let w_i denote the fraction of wealth invested in the i^{th} security at time $t = 0$. Note that we must have $\sum_{i=1}^n w_i = 1$ for any portfolio. Consider now two portfolios:

Portfolio A: 100% invested in security # 1 so that $w_1 = 1$ and $w_i = 0$ for $i = 2, \dots, n$.

Portfolio B: An equi-weighted portfolio so that $w_i = 1/n$ for $i = 1, \dots, n$.

Let R_A and R_B denote the random returns of portfolios A and B , respectively. We immediately have

$$\begin{aligned} E[R_A] &= E[R_B] = \mu \\ \text{Var}(R_A) &= \sigma^2 \\ \text{Var}(R_B) &= \sigma^2/n. \end{aligned}$$

The two portfolios therefore have the same expected return but very different return variances. A risk-averse investor should clearly prefer portfolio B because this portfolio benefits from diversification without sacrificing any expected return. This was the central insight of Markowitz who (in his framework) recognized that investors seek to minimize variance for a given level of expected return or, equivalently, they seek to maximize expected return for a given constraint on variance.

Before formulating and solving the mean variance problem consider Figure 1 below. There were $n = 8$ securities with given mean returns, variances and covariances. We generated $m = 200$ random portfolios from these n securities and computed the expected return and volatility, i.e. standard deviation, for each of them. They are plotted in the figure and are labelled "inefficient". This is because every one of these random portfolios can be improved. In particular, for the same expected return it is possible to find an "efficient" portfolio with a smaller volatility. Alternatively, for the same volatility it is possible to find an efficient portfolio with higher expected return.

1.1 The Efficient Frontier without a Risk-free Asset

We will consider first the mean-variance problem when a risk-free security is not available. We assume that there are n risky securities with the corresponding return vector, \mathbf{R} , satisfying

$$\mathbf{R} \sim \text{MVN}_n(\mu, \Sigma). \quad (1)$$




Figure 1: Sample Portfolios and the Efficient Frontier (without a Riskfree Security).

The mean-variance portfolio optimization problem is formulated as:

$$\begin{aligned} \min_{\mathbf{w}} \quad & \frac{1}{2} \mathbf{w}' \Sigma \mathbf{w} \\ \text{subject to} \quad & \mathbf{w}' \boldsymbol{\mu} = p \\ \text{and} \quad & \mathbf{w}' \mathbf{1} = 1. \end{aligned} \tag{2}$$

Note that the specific value of p will depend on the risk aversion of the investor. This is a simple quadratic optimization problem and it can be solved via standard Lagrange multiplier methods.

Exercise 1 Solve the mean-variance optimization problem (2).

When we plot the mean portfolio return, p , against the corresponding minimal portfolio volatility / standard deviation we obtain the so-called *portfolio frontier*. We can also identify the portfolio having minimal variance among all risky portfolios: this is called the *minimum variance portfolio*. The points on the portfolio frontier with expected returns greater than the minimum variance portfolio's expected return, \bar{R}_{mv} say, are said to lie on the *efficient frontier*. The efficient frontier is plotted as the upper blue curve in Figure 1 or alternatively, the blue curve in Figure 2.

Exercise 2 Let \mathbf{w}_1 and \mathbf{w}_2 be (mean-variance) efficient portfolios corresponding to expected returns r_1 and r_2 , respectively, with $r_1 \neq r_2$. **Show that all efficient portfolios can be obtained as linear combinations of \mathbf{w}_1 and \mathbf{w}_2 .**

The result of the previous exercise is sometimes referred to as a **2-fund theorem**.

1.2 The Efficient Frontier with a Risk-free Asset

We now assume that there is a risk-free security available with risk-free rate equal to r_f . Let $\mathbf{w} := (w_1, \dots, w_n)'$ be the vector of portfolio weights on the n risky assets so that $1 - \sum_{i=1}^n w_i$ is the weight on the risk-free security. An investor's portfolio optimization problem may then be formulated as

$$\begin{aligned} \min_{\mathbf{w}} \quad & \frac{1}{2} \mathbf{w}' \Sigma \mathbf{w} \\ \text{subject to} \quad & \left(1 - \sum_{i=1}^n w_i\right) r_f + \mathbf{w}' \boldsymbol{\mu} = p. \end{aligned} \tag{3}$$




Figure 2: The Efficient Frontier with a Riskfree Security.

The optimal solution to (3) is given by

$$\mathbf{w} = \xi \Sigma^{-1}(\mu - r_f \mathbf{1}) \quad (4)$$

where $\xi := \sigma_{min}^2 / (p - r_f)$ and σ_{min}^2 is the minimized variance, i.e., twice the value of the optimal objective function in (3). It satisfies

$$\sigma_{min}^2 = \frac{(p - r_f)^2}{(\mu - r_f \mathbf{1})' \Sigma^{-1} (\mu - r_f \mathbf{1})} \quad (5)$$

where $\mathbf{1}$ is an $n \times 1$ vector of ones. While ξ (or p) depends on the investor's level of risk aversion it is often inferred from the market portfolio. For example, if we take $p - r_f$ to denote the average excess market return and σ_{min}^2 to denote the variance of the market return, then we can take $\sigma_{min}^2 / (p - r_f)$ as the average or market value of ξ .

Suppose now that $r_f < \bar{R}_{mv}$. When we allow our portfolio to include the risk-free security the efficient frontier becomes a straight line that is tangential to the risky efficient frontier and with a y -intercept equal to the risk-free rate. This is plotted as the red line in Figure 2. That the efficient frontier is a straight line when we include the risk-free asset is also clear from (5) where we see that σ_{min} is linear in p . Note that this result is a **1-fund theorem** since every investor will optimally choose to invest in a combination of the risk-free security and a single risky portfolio, i.e. the **tangency portfolio**. The tangency portfolio, \mathbf{w}^* , is given by the optimal \mathbf{w} of (4) except that it must be scaled so that its component sum to 1. (This scaled portfolio will not depend on p .)

Exercise 3 Without using (5) show that the efficient frontier is indeed a straight line as described above. Hint: consider forming a portfolio of the risk-free security with any risky security or risky portfolio. Show that the mean and standard deviation of the portfolio varies linearly with α where α is the weight on the risk-free-security. The conclusion should now be clear.

Exercise 4 Describe the efficient frontier if no borrowing is allowed.

The **Sharpe ratio** of a portfolio (or security) is the ratio of the expected excess return of the portfolio to the portfolio's volatility. The **Sharpe optimal portfolio** is the portfolio with maximum Sharpe ratio. It is straightforward to see in our mean-variance framework (with a risk-free security) that the tangency portfolio, \mathbf{w}^* , is the Sharpe optimal portfolio.

1.3 Including Portfolio Constraints

We can easily include linear portfolio constraints in the problem formulation and still easily solve the resulting quadratic program. No-borrowing or no short-sales constraints are examples of linear constraints as are leverage

and sector constraints. While analytic solutions are generally no longer available, the resulting problems are still easy to solve numerically. In particular, we can still determine the efficient frontier.

1.4 Weaknesses of Traditional Mean-Variance Analysis

The traditional mean-variance analysis of Markowitz has many weaknesses when applied naively in practice. They include:

1. The tendency to produce extreme portfolios combining extreme shorts with extreme longs. As a result, portfolio managers generally do not trust these extreme weights. This problem is typically caused by estimation errors in the mean return vector and covariance matrix.




Figure 3: The Efficient Frontier, Estimated Frontiers and Realized Frontiers.

Consider Figure 3, for example, where we have plotted the same efficient frontier (of risky securities) as in Figure 2. In practice, investors can never compute this frontier since they do not know the true mean vector and covariance matrix of returns. The best we can hope to do is to approximate it. But how might we do this? One approach would be to simply estimate the mean vector and covariance matrix using historical data. Each of the black dashed curves in Figure 3 is an *estimated* frontier that we computed by: (i) simulating $m = 24$ sample returns from the true (in this case, multivariate normal) distribution (ii) estimating the mean vector and covariance matrix from this simulated data and (iii) using these estimates to generate the (estimated) frontier. Note that the blue curve in Figure 3 is the true frontier computed using the true mean vector and covariance matrix.

The first observation is that the estimated frontiers are quite random and can differ greatly from the true frontier. They may lie below or above the true frontier or they may cross it and an investor who uses such an estimated frontier to make investment decisions may end up choosing a poor portfolio. How poor? The dashed red curves in Figure 3 are the *realized* frontiers that depict the true portfolio mean - volatility tradeoff that results from making decisions based on the estimated frontiers. In contrast to the estimated frontiers, the realized frontiers must always (why?) lie below the true frontier. In Figure 3 some of the realized frontiers lie very close to the true frontier and so in these cases an investor would do very well. But in other cases the realized frontier is far from the (generally unobtainable) true efficient frontier.

These examples serve to highlight the importance of estimation errors in any asset allocation procedure. Note also that if we had assumed a heavy-tailed distribution for the true distribution of portfolio returns then we might expect to see an even greater variety of sample mean-standard deviation frontiers. In addition, it is worth emphasizing that in practice we may not have as many as 24 relevant observations available. For example, if our data observations are weekly returns, then using 24 of them to estimate the joint distribution of returns is hardly a good idea since we are generally more concerned with estimating

conditional return distributions and so **more weight should be given to more recent returns**. A more sophisticated estimation approach should therefore be used in practice. More generally, **it must be stated that estimating expected returns using historical data is very problematic and is not advisable!**

2. The portfolio weights tend to be extremely sensitive to very small changes in the expected returns. For example, even a small increase in the expected return of just one asset can dramatically alter the optimal composition of the entire portfolio. Indeed let \mathbf{w} and $\hat{\mathbf{w}}$ denote the true optimal and estimated optimal portfolios, respectively, corresponding to the true mean return vector, μ , and the sample mean return vector, $\hat{\mu}$, respectively. Then Best and Grauer (1991) showed that

$$\|\mathbf{w} - \hat{\mathbf{w}}\| \leq \xi \|\mu - \hat{\mu}\| \frac{1}{\gamma_{\min}} \left(1 + \frac{\gamma_{\max}}{\gamma_{\min}} \right)$$

where γ_{\max} and γ_{\min} are the largest and smallest eigen values, respectively, of the covariance matrix, Σ . Therefore the sensitivity of the portfolio weights to errors in the mean return vector grows as the ratio $\gamma_{\max}/\gamma_{\min}$ grows. But this ratio, when applied to the estimated covariance matrix, $\hat{\Sigma}$, typically becomes large as the number of asset increases and the number of sample observations is held fixed. As a result, we can expect large errors for large portfolios with relatively few observations.

3. While it is commonly believed that errors in the estimated means are of much greater significance, errors in estimated covariance matrices can also have considerable impact. While it is generally easier to estimate covariances than means, the presence of heavy tails in the return distributions can result in significant errors in covariance estimates as well. These problems can be mitigated to varying extents through **the use of more robust estimation techniques**.

As a result of these weaknesses, portfolio managers traditionally have had little confidence in mean-variance analysis and therefore applied it very rarely in practice. Efforts to overcome these problems include the use of better estimation techniques such as **the use of shrinkage estimators, robust estimators and Bayesian techniques such as the Black-Litterman framework introduced in the early 1990's**. (In addition to mitigating the problem of extreme portfolios, the Black-Litterman framework allows users to specify their own subjective views on the market in a consistent and tractable manner.) Many of these techniques are now used routinely in general asset allocation settings. It is worth mentioning that the problem of extreme portfolios can also be mitigated in part by placing **no short-sales and / or no-borrowing constraints** on the portfolio.




Figure 4: Robust Estimation of the Efficient Frontier.

In Figure 4 above we have shown an estimated frontier that was computed using a more robust estimation procedure. We see that it lies much closer to the true frontier which is also the case with it's corresponding realized frontier.

2 The Capital Asset Pricing Model (CAPM)

If every investor is a mean-variance optimizer then we can see from Figure 2 and our earlier discussion that each of them will hold the same tangency portfolio of risky securities in conjunction with a position in the risk-free asset. **Because the tangency portfolio is held by all investors and because markets must clear**, we can identify this portfolio as the **market portfolio**. **The efficient frontier is then termed the capital market line**.

Now let R_m and \bar{R}_m denote the return and expected return, respectively, of the market, i.e. tangency, portfolio. The central insight of the *Capital Asset-Pricing Model* is that in equilibrium the riskiness of an asset is not measured by the standard deviation of its return but by its **beta**. In particular, there is a linear relationship between the expected return, $\bar{R} = E[R]$ say, of any security (or portfolio) and the expected return of the market portfolio. It is given by

$$\bar{R} = r_f + \beta (\bar{R}_m - r_f) \quad (6)$$

where $\beta := \text{Cov}(R, R_m) / \text{Var}(R_m)$. In order to prove (6), consider a portfolio with weights α and weight $1 - \alpha$ on the risky security and market portfolio, respectively. Let R_α denote the (random) return of this portfolio as a function of α . We then have

$$E[R_\alpha] = \alpha \bar{R} + (1 - \alpha) \bar{R}_m \quad (7)$$

$$\sigma_{R_\alpha}^2 = \alpha^2 \sigma_R^2 + (1 - \alpha)^2 \sigma_{R_m}^2 + 2\alpha(1 - \alpha) \sigma_{R, R_m} \quad (8)$$

where $\sigma_{R_\alpha}^2$, σ_R^2 and $\sigma_{R_m}^2$ are the returns variances of the portfolio, security and market portfolio, respectively. We use σ_{R, R_m} to denote $\text{Cov}(R, R_m)$. Now note that as α varies, the mean and standard deviation, $(E[R_\alpha], \sigma_{R_\alpha}^2)$, trace out a curve **that cannot (why?) cross the efficient frontier**. This curve is depicted as the dashed curve in Figure 5 below.




Figure 5: Proving the CAPM relationship.

Therefore at $\alpha = 0$ this curve must be tangent to the capital market line. Therefore the slope of the curve at

$\alpha = 0$ must equal the slope of the capital market line. Using (7) and (8) we see the former slope is given by

$$\begin{aligned}
 \left. \frac{dE[R_\alpha]}{d\sigma_{R_\alpha}} \right|_{\alpha=0} &= \left. \frac{dE[R_\alpha]}{d\alpha} \right/ \left. \frac{d\sigma_{R_\alpha}}{d\alpha} \right|_{\alpha=0} \\
 &= \left. \frac{\sigma_{R_\alpha} (\bar{R} - \bar{R}_m)}{\alpha\sigma_R^2 - (1-\alpha)\sigma_{R_m}^2 + (1-2\alpha)\sigma_{R,R_m}} \right|_{\alpha=0} \\
 &= \frac{\sigma_{R_m} (\bar{R} - \bar{R}_m)}{-\sigma_{R_m}^2 + \sigma_{R,R_m}}.
 \end{aligned} \tag{9}$$

The slope of the capital market line is $(\bar{R}_m - r_f) / \sigma_{R_m}$ and equating the two therefore yields

$$\frac{\sigma_{R_m} (\bar{R} - \bar{R}_m)}{-\sigma_{R_m}^2 + \sigma_{R,R_m}} = \frac{\bar{R}_m - r_f}{\sigma_{R_m}} \tag{10}$$

which upon simplification gives (6).

The CAPM result is one of the most famous results in all of finance and, even though it arises from a simple one-period model, it provides considerable insight to the problem of asset-pricing. For example, it is well-known that riskier securities should have higher expected returns in order to compensate investors for holding them. But how do we measure risk? Counter to the prevailing wisdom at the time the CAPM was developed, the riskiness of a security is not measured by its return volatility. Instead it is measured by its beta which is proportional to its covariance with the market portfolio. This is a very important insight. Nor, it should be noted, does this contradict the mean-variance formulation of Markowitz where investors do care about return variance. Indeed, we derived the CAPM from mean-variance analysis!

Exercise 5 Why does the CAPM result not contradict the mean-variance problem formulation where investors do measure a portfolio's risk by its variance?

The CAPM is an example of a so-called 1-factor model with the market return playing the role of the single factor. Other factor models can have more than one factor. For example, the Fama-French model has three factors, one of which is the market return. Many empirical research papers have been written to test the CAPM. Such papers usually perform regressions of the form

$$R_i - r_f = \alpha_i + \beta_i (R_m - r_f) + \epsilon_i \tag{11}$$

where α_i (not to be confused with the α we used in the proof of (6)) is the intercept and ϵ_i is the *idiosyncratic* or residual risk which is assumed to be independent of R_m and the idiosyncratic risk of other securities. If the CAPM holds then we should be able to reject the hypothesis that $\alpha_i \neq 0$. The evidence in favor of the CAPM is mixed. But the language inspired by the CAPM is now found throughout finance. For example, we use β 's to denote factor loadings and α 's to denote excess returns even in non-CAPM settings.